


Figure 6: Multi-Model VLM Evaluation - Executive Summary

BEST MODEL: GEMINI 3 PRO

Accuracy: 95.5% | IoU: 95.5% | Anomalies: 3 | Time: 344s



KEY FINDINGS:

- ✓ gemini-2.5-pro achieves 93.75% accuracy - closest to human performance (98.25%)
 - ✓ 12% improvement over Flash baseline with 56% fewer anomalies
- ✓ Gemini-3-pro-preview underperforms (90.48%) - likely due to preview status
- ✓ All models struggle with amb_d/to_drawer_close and le_r_block (spatial reasoning)
 - ✓ prefix_frames mode successfully integrates temporal information

RECOMMENDATION: Deploy gemini-2.5-pro for production VLM-based goal inference