


VLM Goal Inference Pipeline



- Key Features:
- Cumulative frame input (prefix frames $f_0 \dots f_i$)
 - Vision-Language Model: Google Gemini
 - Multi-modal reasoning (vision + text)
 - Real-time goal inference capability

A: Move to left block / Close bottom drawer

B: Move to right block / Close top drawer

C: Uncertain / Insufficient information

VLM Performance: 64.5% accuracy | 74.2% human agreement